

# Analyzing RALMs with Selective State Space and Attention based Architectures for Long Sequence Modeling

Sebastian P. Jaskowski<sup>1</sup>

Nolan M. Bridges<sup>1</sup>

Austin T. Barton<sup>1</sup>

<sup>1</sup>Georgia Institute of Technology

sjaskowski3@gatech.edu

bridges@gatech.edu

abarton40@gatech.edu

## 1 Introduction

Recently, a new language model paradigm has emerged in the form of Retrieval Augmented Language Models, or RALMs (Lewis et al., 2021; Gao et al., 2024). These models improve upon stand-alone language models by utilizing a retriever to acquire and inject relevant context information into a language model prompt during inference time. By injecting knowledge into the model at runtime, this paradigm addresses limitations such as outdated parameterized knowledge, the generation of false or incorrect statements, and the lack of supportive reasoning for responses.

Most recent language models, and thus almost all Retrieval Augmented Language Models, have been built upon the transformer architecture. (Vaswani et al., 2023). While the transformer has allowed for incredibly powerful language models due to its ability to selectively learn relationships between arbitrary elements in sequential data, it also introduces drawbacks such as quadratic scaling over sequence length, resulting in limited context windows (Xu et al., 2024), and difficulty considering all information within its context window even when the information fits. (Liu et al., 2023). These limitations of transformer-based models are particularly impactful on RALMs, as they effectively place an upper-limit on the number of retrieved context chunks a RALM can process.

On a different note, recent work in the field of state space models (SSMs), most notably in the form of the newly proposed Mamba architecture, has shown strong results on sequence modelling tasks, especially on long sequences with Long Range Dependencies, or LRDs (Gu et al., 2022, 2021). This new architecture holds promising potential for its use as a language model for RALMs, due to its ability to process and utilize a higher number of retrieved chunks, potentially leading to higher accuracy on QA tasks. Specifically, it




Figure 1: Scatter plot of accuracy results of Mamba-Chat-2.8B and Dolly-v2-3B over each  $k$  value with  $B$ -Spline ( $B = 4$ ) interpolation.

may see higher performance when there are more retrieved chunks to parse through, due to its demonstrated stronger ability at signal/noise disambiguation.

As such, we aim to assess these claims by providing the first comprehensive, end-to-end evaluation and benchmarking of a RALM constructed upon a language model based on the Mamba architecture. We construct a RALM based on the 2.8B parameter instruction-tuned Mamba-Chat model (Mattern and Hohn, 2023) and evaluate it on a subset of the popular QA benchmarking dataset TriviaQA (Joshi et al., 2017). We compare our results to the performance of a RALM based on the 2.8B parameter instruction-tuned Dolly-v2-3b model (Conover et al., 2023), and we evaluate both models over  $k$  values ranging from 1 to 10 (where  $k$  refers to the number of context chunks retrieved from the vector database and fed to the language model).

We find that both models see similar performance for small to medium values of  $k$  ( $k \leq 7$ ), and the performance of both models declines as  $k$  grows. This contradicts our initial assumption that the Mamba-Chat model would see a signifi-

cant, overall performance increase when compared to Dolly-v2-3b. However, when large amounts of context are presented ( $k > 7$ ), we see that the performance degradation faced by Dolly-v2-3b is much greater than that faced by the Mamba-Chat model. This suggests that Mamba-Chat may be more resilient to the larger amounts of irrelevant information present when  $k$  is larger.

## 2 Background

### 2.1 Retrieval-Augmented Language Models

For the past few years, Large Language Models (LLMs) have dominated the field of Natural Language Processing (NLP) due to their ability to function as Few-Shot Learners (Brown et al., 2020), achieving state of the art performance on a variety of NLP tasks without task-specific training. However, key limitations remain, such as their inability to update their parameterized knowledge base to incorporate new information, their tendency to produce made-up information (a behavior that has been referred to as "hallucinations"), and their inability to source their information, which reduces user trust in the model. (Lewis et al., 2021; Huang et al., 2023).

As a result, RALMs emerged as a effective solution to address these drawbacks by providing an external non-parametric knowledge base in the form of a vector index, from which a the model can retrieve relevant knowledge context that can then be appended to the initial prompt (Lewis et al., 2021). By having relevant knowledge in the prompt, the LLM no longer relies on its parametric knowledge, which is often inaccurate and leads to the common problems described above.

Retrieval-augmented models operate in a multi-step process. Before inference, a large corpus of text that contains potentially relevant information to future prompts is chunked, embedded using an embedding model, and stored in a vector database that maps vector representations to their corresponding chunks. At inference time, the query to the model is embedded using the same embedding model as used during database creation, and the  $k$  most relevant chunks (where  $k$  is a hyperparameter), as determined by cosine similarity of embeddings, is retrieved. These chunks are then combined with the original query to form a language model prompt, which is passed to the LLM.

While RALMs have found tremendous success in QA tasks, recent studies have highlighted that

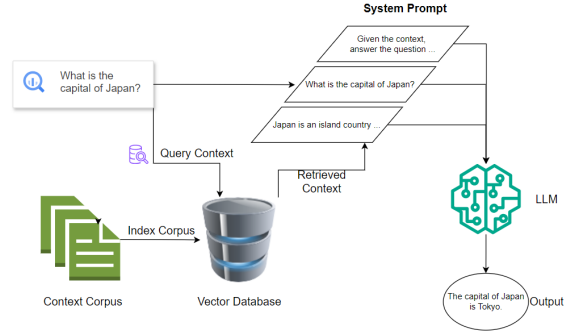


Figure 2: Diagram of standard Retrieval-Augmented Language Model

transformer-based Retrieval-Augmented Language Models (RALMs) face significant challenges when handling extensive amounts of retrieved knowledge. The inherent quadratic memory constraints of the self-attention mechanism limit the context window, which in turn poses difficulties in integrating sufficient context. This limitation has spurred investigations into various information condensation strategies, including PRCA (Yang et al., 2023), RECOMP (Xu et al., 2023), and the Filter-Reranker approach (Ma et al., 2023). Furthermore, evidence from recent research suggests that even as the context windows are enlarged in more substantial models, transformer language models exhibit limitations in effectively leveraging all pertinent content chunks (Xu et al., 2024) (Liu et al., 2023).

### 2.2 State Space Models

These challenges have a potential solution in State Space Models, or SSMs. While research into SSMs fell out of the limelight due to the prominent success of the transformer architectures, State Space Models promise a key significant improvement over transformers in the form of linear scaling over sequence length, as opposed to the quadratic scaling found in transformers. In 2021, researchers developed the Linear State Space Layer (LSSL), a model that combines linear recurrence with a continuous-time state space (Gu et al., 2021). This model was further improved upon by Structured State Space Models, commonly referred to as S4, which alleviated the heavy computational requirements that previously hindered LSSLs (Gu et al., 2022). These models saw great performance improvements over previous state-of-the-art, yet still struggled to match the performance and efficiency of Transformers on language tasks (Gu and Dao, 2023).

More recently, the Mamba architecture emerged. This variation of the S4 model saw an added selectivity mechanism, allowing the model to become context-aware and selectively consider specific sequence elements while maintaining fast computation and efficient memory usage. This improvement ended up being crucial to boosting the performance of State Space Models; Mamba achieved state-of-the-art performance on sequences up to 1M tokens, and 5 times faster inference than similarly sized transformers (Gu and Dao, 2023).

### 3 Methodology

We construct our RALM models using a common Retrieval-Augmented Generation (RAG) framework for all language model architectures, ensuring that the only difference between RALMs is the underlying language model. All models utilize the same retriever, vector index, and prompting format.

#### 3.1 Vector Index

To support efficient and accurate context retrieval, we constructed our vector database using the Facebook AI Similarity Search (FAISS) vector database (Johnson et al., 2017). We began by aggregating 10,031 wikipedia articles designated as ground truth context in the TriviaQA dataset (Joshi et al., 2017). This set of articles is a subset of a larger context base of 73,976 wikipedia articles, however we removed all articles that were not designated as a ground truth reference by any question in our QA dataset in order to alleviate the prohibitively expensive computational requirements of embedding such a high amount of text.

These 10,031 wikipedia articles, totalling 223MB of data, were then chunked (with chunk size equal to 1000 characters) and passed through the *bge-small-en-v1.5* embedding model (Xiao et al., 2023) to create vector representations of chunks. Context embedding remains a computational bottleneck for RALMs; thus, our choices of chunk size and embedding model reflect this limitation. We chose the *bge-small-en-v1.5* embedding model due to its small size of only 0.13GB and relatively high performance. We set our target chunk size at 1000<sup>1</sup> as smaller chunk sizes would have resulted in prohibitively expensive chunking and embedding time.

<sup>1</sup>While we set our target chunk size to 1000, chunk lengths vary due to long sequences of consecutive characters without the occurrence of our split character (a newline character).

Model	Mamba-Chat	Dolly-v2-3b
Num. Parameters	2.8B	2.8B
Context Limit	unlimited	2048 tokens
Instruction-tuned?	Yes	Yes
Temperature	0	0

Table 1: Comparison of Models Evaluated

Upon completion of indexing, the original 10,031 wikipedia articles resulted in 245,624 chunks stored in our vector database.

#### 3.2 Retriever

To perform retrieval, we utilized the FAISS library to perform efficient k-Nearest Neighbors search based on cosine similarity over our vector database. We chose to utilize the FAISS library due to its efficient utilization of GPUs when performing similarity search, which significantly sped up RALM inference (Johnson et al., 2017).

#### 3.3 Prompt Formatting

To ensure maximum fairness in evaluation, we standardized our prompt across all models evaluated. Our prompt was constructed based on careful review of previously used system prompts for Retrieval-Augmented Language Models (Azhar, 2024).

Our prompt is displayed below.

*Please respond to the original query. If the selected document context is relevant and informative, provide a detailed answer based on its content. However, if the selected document context does not offer useful information or is not applicable, simply state 'No answer found'*

#### 3.4 Language Models

We evaluate two language models within our Retrieval-Augmented Generation structure: the 2.8B parameter instruction-tuned Mamba-Chat model based on the Mamba architecture and the 2.8B parameter instruction-tuned Dolly-v2-3b model based on the transformer architecture.

We chose the Mamba-Chat model (Mattern and Hohr, 2023) as it was one of few existing instruction-tuned language models based on the Mamba architecture. We then chose Dolly-v2-3b (Conover et al., 2023) as our baseline transformer

model due to its relatively-decent success, and comparable model size to Mamba-Chat.

### 3.5 Inference and Evaluation Methodology

In order to assess the ability of our RALMs, we evaluate the models over a subset of the popular Question-Answering (QA) benchmark TriviaQA (Joshi et al., 2017). This dataset features a set of questions, with each question matched to a particular noun answer, as well as a list of possible aliases for that answer.

Due to an average model inference time in excess of 10 seconds, evaluating over the entire dataset is beyond the capabilities of our computational resources. As such, we chose to evaluate our models over a random sample of 1,000 questions from the TriviaQA dataset.

Both models were evaluated with varying amounts of retrieved context, with  $k$  ranging from 1 to 10. An upper-bound at 10 was chosen to due increasing evaluation times at higher values of  $k$  (due to greater execution time of Nearest Neighbor Search in retrieval), as well as context window limitations on Dolly-v2-3b. We then perform an additional evaluation on Mamba-Chat at  $k = 20$ , yet this result is incomparable since  $k = 20$  would far exceed the context length on Dolly-v2-3B.

To evaluate whether a given model-generated answer is correct, we checked if any alias of the correct answer appeared in the model output text. We refer to this evaluation approach as Exact Subsequence Matching. Due to the nature of this dataset featuring answers that are proper nouns, we believed that this evaluation methodology accurately reflects the performance of this model in this particular QA task.

Evaluations were performed on a High-RAM Google Colab instance utilizing a T4 GPU.

## 4 Results

We find that both Mamba-Chat and Dolly-v2-3b perform similarly for small to medium amounts of retrieved context, and while both models encounter performance degradation as  $k$  increases, the Mamba-Chat model remains more resilient to these declines in performance. Our results are summarized in Table 2 and visualized in Figure 1.

For small to medium values of  $k$ , we find that both models achieve similar performance, with Dolly-v2-3b slightly outperforming its counterpart. This contradicts our initial assumption that Mamba-

Model	Mamba-Chat	Dolly-v2-3b
k=1	44.6%	43.2%
k=2	41.6%	44.9%
k=3	43.2%	45.8%
k=4	44.5%	43.9%
k=5	44.0%	46.3%
k=6	43.8%	47.0%
k=7	44.0%	47.8%
k=8	38.8%	43.6%
k=9	35.2%	29.8%
k=10	29.6%	10.7%
k=20	21.5%	N/A

Table 2: Accuracy of Models on TriviaQA dataset-subset.

Chat would see overall greater performance. We acknowledge that the performance for small and medium values of  $k$  is likely bottle-necked by the retrieval process, and the use of a more powerful embedding model may lead to higher performance overall, as well as the occurrence of a performance disparity at higher accuracy levels.

As  $k$  grows to larger values, we see that the performance of both models begins to decline, with the performance of Dolly-v2-3B declining significantly greater than that of Mamba-Chat. At  $k = 10$ , we find that Mamba-Chat achieves an Exact Subsequence Match accuracy of 29.6%, while Dolly-v2-3B achieves only 10.7%.

We infer that this performance decline is caused by two factors: (1) context window limitations in the transformer-based model and (2) model’s inability to effectively consider all context in its input.

The former of these is only applicable to Dolly-v2-3b, as the nature of the Mamba architecture does not necessitate an explicit context length. Our maximum  $k$  value of 10 was chosen such that model input did not exceed 2048 tokens<sup>2</sup>, but due to variations in tokenization and chunking, we cannot guarantee that prompts did not exceed the context window length of Dolly-v2-3b. This leads us to conclude that context window length limitations may have been a factor in the sudden decrease in performance of Dolly-v2-3b at larger values of  $k$ .

The second potential factor for performance decreases in both models is the ability of the model to properly utilize all inputted context chunks. Due

<sup>2</sup>This is based on an assumption of 5 characters per token, resulting in 200 tokens per context chunk, which means 2000 tokens of total context at  $k = 10$



to previous work demonstrating the strong performance of the Mamba architecture on modeling Long Range Dependencies (Gu and Dao, 2023), we believed that Mamba-Chat would continue to be performant as inputted context lengths increased. Our results show that while performance reduction in Mamba-Chat as  $k \rightarrow 10$  does occur, it is not as significant as that seen in Dolly-v2-3b. This leads us to confirm our initial hypothesis: Mamba-Chat does demonstrate stronger resilience to increased context input than transformer-based models. This is particularly evident with Mamba-Chat maintaining an Exact Subsequence Match accuracy of 21.5% at  $k = 20$ , while the accuracy of Dolly-v2-3b falls to 10.7% by  $k = 10$ .

## 5 Conclusion

In this work, we aimed to perform a comprehensive, end-to-end evaluation of the Mamba architecture as a language model for RALM systems, comparing it to the traditional transformer-language model. We evaluate both models on a long-context QA task, utilizing the popular TriviaQA QA evaluation dataset. We find that the Mamba-Chat model does not perform better than a traditional transformer-based language model at low to medium amounts of retrieved context. However, it maintains significantly higher performance, as compared to Dolly-v2-3b, when the amount of retrieved context is high. While both models suffer from performance degradation, Mamba-Chat appears to be significantly more resilient to this degradation.

## 6 Limitations and Future Work

While we aimed to provide a full comprehensive evaluation of differences between our two RALMs, our limited computational resources resulted in some limitations, most notably in the following areas:

1. **Small Number of Evaluated Models.** Due to computational challenges and time constraints, we were only able to compare the Mamba-based RALM to one Transformer-based RALM. There is significant room for future work in comparing the Mamba RALM model to other transformer models, including even models with larger numbers of parameters such 7B parameter models and 13B parameter models, as well as models with larger context length limitations, allowing for evaluation on higher values of  $k$ .
2. **Limited Evaluation Dataset.** Once again due to computational resource constraints, we were only able to benchmark the models on a small subset of the TriviaQA dataset. In future work, we would be interested in performing a more comprehensive evaluation against the entire dataset, as well as expanding the evaluation to other prominent QA benchmarks such as Natural Questions QA (Kwiatkowski et al., 2019).
3. **Small Encoding Model** While the embedding model we chose, *bge-small-en-v1.5*, is a high performing model given its size, there are higher performing embedding models available. Unfortunately, we were unable to utilize them due to computational constraints, as larger models would have required more memory usage and embedding time.
4. **Truncated Model Parameters** Due to computational limitations with regards to GPU memory, we were not able to use the models' parameters as they were trained, using 16-bit float parameters instead of the default 32 bits. This may have induced slight errors on larger sequences, but we believe that this effect was likely negligible.

In summary, future work in this direction could focus on expanding our evaluation to more transformer-based models, as well as expanding the evaluation to new QA datasets. In particular, future work that attempts to expand our evaluations to larger transformer models with longer context lengths may significantly strengthen the significance of our results.

## References

- Azhar. 2024. Implementing rag with mamba and the qdrant database: A detailed exploration (with code).
- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. [Language models are few-shot learners](#).
- Mike Conover, Matt Hayes, Ankit Mathur, Jianwei Xie, Jun Wan, Sam Shah, Ali Ghodsi, Patrick Wendell, Matei Zaharia, and Reynold Xin. 2023. [Free dolly: Introducing the world’s first truly open instruction-tuned llm](#).
- Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Qianyu Guo, Meng Wang, and Haofen Wang. 2024. [Retrieval-augmented generation for large language models: A survey](#).
- Albert Gu and Tri Dao. 2023. [Mamba: Linear-time sequence modeling with selective state spaces](#).
- Albert Gu, Karan Goel, and Christopher Ré. 2022. [Efficiently modeling long sequences with structured state spaces](#).
- Albert Gu, Isys Johnson, Karan Goel, Khaled Saab, Tri Dao, Atri Rudra, and Christopher Ré. 2021. [Combining recurrent, convolutional, and continuous-time models with linear state-space layers](#).
- Lei Huang, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, and Ting Liu. 2023. [A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions](#).
- Jeff Johnson, Matthijs Douze, and Hervé Jégou. 2017. [Billion-scale similarity search with gpus](#).
- Mandar Joshi, Eunsol Choi, Daniel S. Weld, and Luke Zettlemoyer. 2017. [Triviaqa: A large scale distantly supervised challenge dataset for reading comprehension](#). *ArXiv*, abs/1705.03551.
- Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, Kristina Toutanova, Llion Jones, Matthew Kelcey, Ming-Wei Chang, Andrew M. Dai, Jakob Uszkoreit, Quoc Le, and Slav Petrov. 2019. [Natural questions: A benchmark for question answering research](#). *Transactions of the Association for Computational Linguistics*, 7:452–466.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2021. [Retrieval-augmented generation for knowledge-intensive nlp tasks](#).
- Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. 2023. [Lost in the middle: How language models use long contexts](#).
- Yubo Ma, Yixin Cao, YongChing Hong, and Aixin Sun. 2023. [Large language model is not a good few-shot information extractor, but a good reranker for hard samples!](#)
- Justus Mattern and Konstantin Hohn. 2023. [Mamba-chat](#). GitHub.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. 2023. [Attention is all you need](#).
- Shitao Xiao, Zheng Liu, Peitian Zhang, and Niklas Muennighoff. 2023. [C-pack: Packaged resources to advance general chinese embedding](#).
- Fangyuan Xu, Weijia Shi, and Eunsol Choi. 2023. [Re-comp: Improving retrieval-augmented lms with compression and selective augmentation](#).
- Peng Xu, Wei Ping, Xianchao Wu, Lawrence McAfee, Chen Zhu, Zihan Liu, Sandeep Subramanian, Evelina Bakhturina, Mohammad Shoeybi, and Bryan Catanzaro. 2024. [Retrieval meets long context large language models](#).
- Haoyan Yang, Zhitao Li, Yong Zhang, Jianzong Wang, Ning Cheng, Ming Li, and Jing Xiao. 2023. [PRCA: Fitting black-box large language models for retrieval question answering via pluggable reward-driven contextual adapter](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 5364–5375, Singapore. Association for Computational Linguistics.

## Appendix A

### A Summary of Individual Contributions

The following table summarizes the specific contributions to the project by individual authors.

Task	Sebastian Jaskowski	Nolan Bridges	Austin Barton
Proposal Writing	25%	10%	65%
Model Code	10%	70%	20%
RAG Framework Code	20%	40%	40%
TriviaQA Dataset Preprocessing	60%	30%	10%
Evidence Embedding/Vector Index Construction	60%	20%	20%
Model Evaluation	45%	10%	45%
Final Report Writing	70%	10%	20%
Presentation Slides	10%	20%	70%